

Turn scattered case files into a **decision-ready** knowledge base



The Problem

Your administrative knowledge — case files, precedent decisions, inquiry responses, handover memos, policy interpretations — is scattered across shared drives and paper archives, and lost with every personnel rotation. Finding how a similar case was decided takes days; consistency depends on whoever remembers.



What Neuron Does

It links every record into one ontology-structured database anchored to each case number, so your team gets instant answers to "how was this decided before?" and "which past cases does this interpretation affect?" — without changing how you work today.

What you can do with the centralized database

01

Surface precedents at intake

Flag incoming cases that match past complex or contested cases, so the right precedent and the right caution are applied from day one.

02

Check impact across related cases

When a rule interpretation changes, instantly see every past and pending case decided under the old reading and judge the full scope of review needed.

03

Verify measures actually achieved outcomes

Track whether programs and administrative measures produced their intended results across similar cases — evidence for audits and budget reviews.

04

Keep decisions consistent through rotation

Surface similar past cases and their rationale at judgment time, so a two-year rotation doesn't reset twenty years of institutional knowledge.

"Every April, rotation used to mean months of relearning. This year, new staff answered precedent questions in their first week."

Division Director - Municipal government

See it with your own data

We start with a focused PoC using your existing records — case files, response logs, handover documents, as they are. No system migration, no process change required.

Let's talk for 30 minutes. →

 the-neuron.com/en

 tharada@the-unchain.com