

Turn scattered incident history into a **decision-ready** knowledge base



The Problem

Your operational knowledge — past incidents, postmortems, runbooks, ticket threads, workarounds — is scattered across Jira, Confluence, Slack, and monitoring tools, and locked in your senior engineers' heads. Finding a similar past incident takes hours; resolution quality depends on who's on call.



What Neuron Does

It links every record into one ontology-structured database anchored to each service and incident ID, so your team gets instant answers to "has this failure pattern happened before?" and "which dependent systems does this affect?" — without changing how you work today.

What you can do with the centralized database

01

Detect incident signs before outages

Flag services showing patterns that preceded past incidents — recurring alerts, config drift, error trends — so risks are caught before users feel them, not after the pager goes off.

02

Check blast radius across dependencies

When one service degrades, instantly see which downstream systems, customers, and SLAs were affected in similar past incidents and judge the full exposure.

03

Verify fixes actually worked

Track whether the same incident class recurred across similar services and releases after each postmortem action item — evidence you can show management and auditors.

04

Resolve at 3 a.m. with full context

Surface similar past incidents, root causes, and proven workarounds at triage time, so on-call engineers act like your most experienced one.

"Root-cause hunting used to mean grepping through old Slack threads and stale wiki pages for hours. Now similar incidents surface in minutes — with the fix that worked."

Head of SRE - B2B software company

See it with your own data

We start with a focused PoC using your existing operational data — tickets, postmortems, runbooks, as they are. No system migration, no process change required.

Let's talk for 30 minutes. →

 the-neuron.com/en

 tharada@the-unchain.com